





```
[16]: seqlist(mydf[['vehicleum']])
# display(seq)

window_size = 15
sliding_window = []
for i in range(len(seq) - window_size + 1):
    sliding_window.append(seq[i: i + window_size])

# print(sliding_window)

sliding_df = pd.DataFrame(sliding_window)

sliding_df = sliding_df.drop(columns=[1,2,3,4,5])
display(sliding_df)
y = sliding_df[0]
x = sliding_df.drop(columns=[0])
x scaler.fit_transform(X)
# print(X')
# display(X')
# print(y')
# display(y')
# X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = .2, random_state=42)
X_train = X[:360]
X_test = X[-360:]
y_train = y[:360]
y_test = y[-360:]
print(X_train.shape)
print(X_test.shape)
print(y_train.shape)
print(y_test.shape)
```

	0	6	7	8	9	10	11	12	13	14
0	329.025281	382.934451	386.343057	346.398775	349.929240	299.717332	318.972804	333.206330	329.415663	304.475108
1	329.865783	386.943057	346.398775	349.929240	299.717332	318.972804	333.206330	329.415663	304.475108	299.694444
2	343.341337	346.398775	349.929240	299.717332	318.972804	333.206330	329.415663	304.475108	299.694444	292.779657
3	344.201131	349.929240	299.717332	318.972804	333.206330	329.415663	304.475108	299.694444	292.779657	296.475399
4	366.609590	299.717332	318.972804	333.206330	329.415663	304.475108	299.694444	292.779657	296.475399	290.586332
...	...	...	...	...	...	...	...	...	...	...
1780	1216.849707	1234.099139	1355.045640	1344.720284	1570.658959	1542.629821	1527.413931	1527.413931	1488.975794	1488.975794
1781	1160.012708	1355.045640	1344.720284	1570.658959	1542.629821	1527.413931	1527.413931	1488.975794	1488.975794	1527.413931
1782	1097.499438	1344.720284	1570.658959	1542.629821	1527.413931	1527.413931	1488.975794	1488.975794	1527.413931	1527.413931
1783	1040.797146	1570.658959	1542.629821	1527.413931	1527.413931	1488.975794	1488.975794	1527.413931	1527.413931	1488.975794
1784	1112.020783	1542.629821	1527.413931	1527.413931	1488.975794	1488.975794	1527.413931	1527.413931	1488.975794	1488.975794

1785 rows × 10 columns

X\_train:

(1425, 9)

X\_test:

(360, 9)

y\_train:

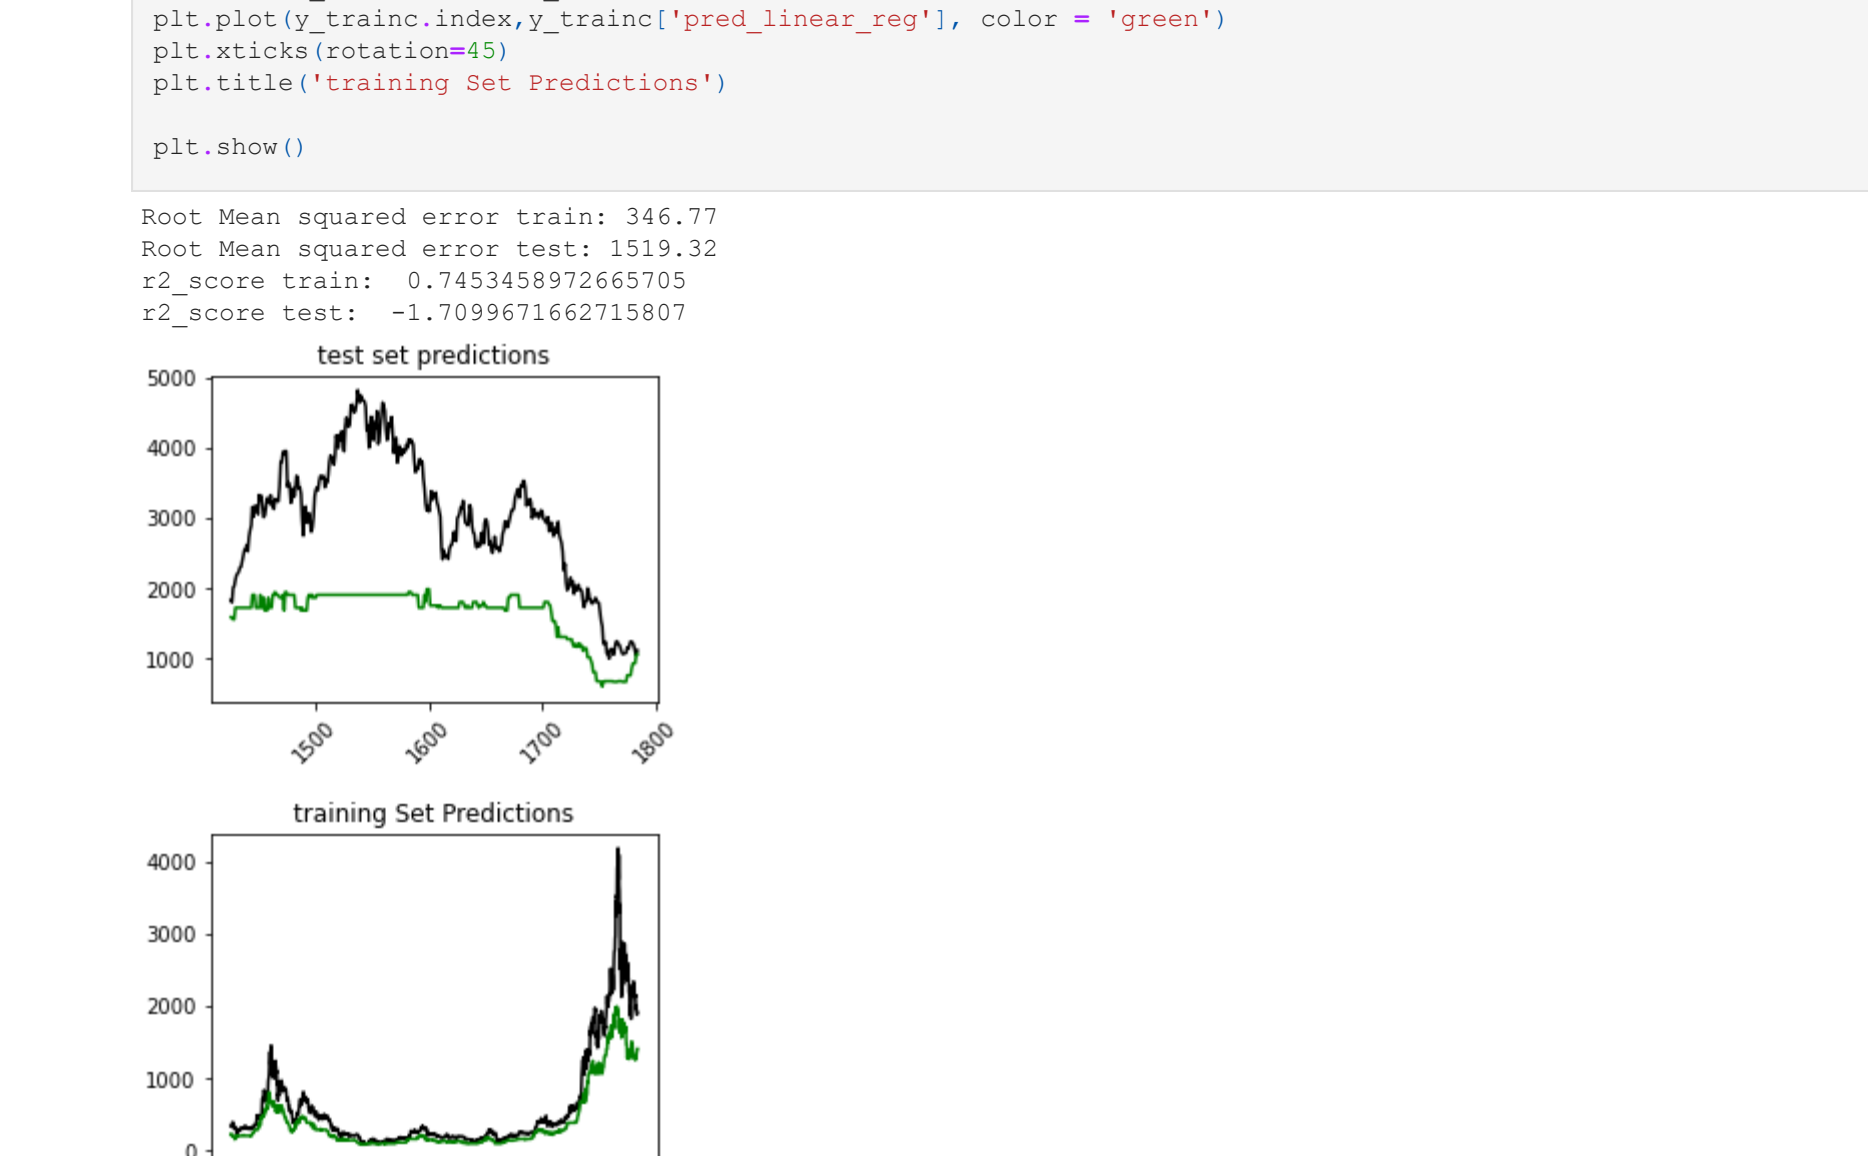
(1425,)

y\_test:

(360,)

```
In [17]: sns.heatmap(sliding_df.corr(), annot=True)
```

Out[17]:



xgboost on sliding window

```
In [18]: xg_reg=xgb.XGBRegressor(objective='reg:squarederror',colsample_bytree=0.3, learning_rate=0.1,
                             max_depth=5,alpha=10,n_estimators=10)
xg_reg.fit(X_train,y_train)
y_pred_test=xg_reg.predict(X_test)
y_pred_train=xg_reg.predict(X_train)
# rmse=np.sqrt(mean_squared_error(y_test,preds))
# print(preds)
# print(y_check)
# print('xgb rmse: ', rmse)

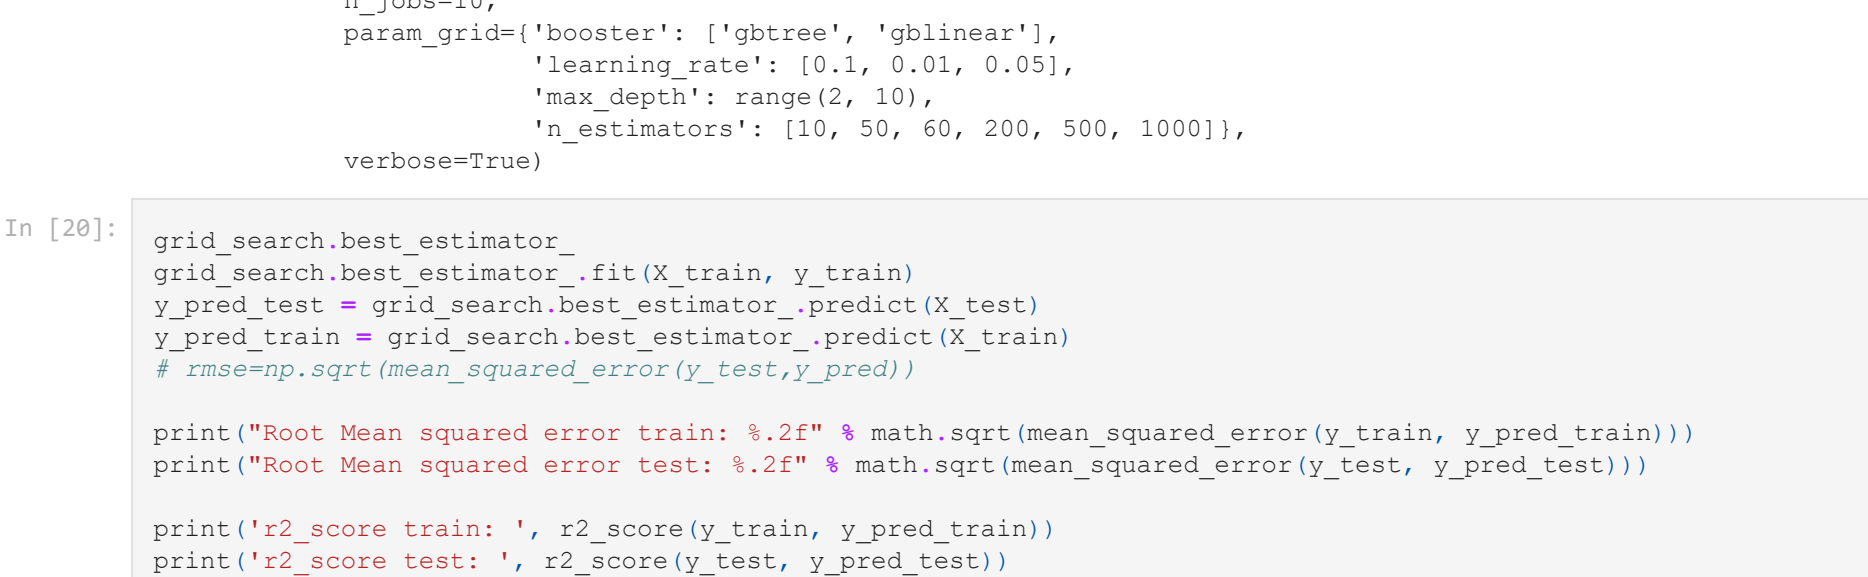
print("Root Mean squared error train: %.2f" % math.sqrt(mean_squared_error(y_train, y_pred_train)))
print("Root Mean squared error test: %.2f" % math.sqrt(mean_squared_error(y_test, y_pred_test)))

print('r2_score train: ', r2_score(y_train, y_pred_train))
print('r2_score test: ', r2_score(y_test, y_pred_test))

y_testc = pd.DataFrame(y_test)
# print(y_check)
y_testc['pred_linear_reg']= y_pred_test
# print(y_check)
figure(figsize=(4, 3))
plt.plot(y_testc.index,y_testc[0], color = 'black')
plt.plot(y_testc.index,y_testc['pred_linear_reg'], color = 'green')
plt.xticks(rotation=45)
plt.title('test set predictions')
plt.show()

y_trainc = pd.DataFrame(y_train)
# print(y_check)
y_trainc['pred_linear_reg']= y_pred_train
# print(y_check)
figure(figsize=(4, 3))
plt.plot(y_trainc.index,y_trainc[0], color = 'black')
plt.plot(y_trainc.index,y_trainc['pred_linear_reg'], color = 'green')
plt.xticks(rotation=45)
plt.title('training Set Predictions')
plt.show()
```

Root Mean squared error train: 346.77  
Root Mean squared error test: 1519.32  
r2\_score train: 0.7453458972646705  
r2\_score test: -1.7099671662715807



## hyperparameter tuning on sliding windows feature set

```
In [19]: parameters = {
    'max_depth': range(2, 10, 1),
    'booster': ['gbtree', 'gblinear'],
    'n_estimators': range(0, 1000, 50),
    'n_estimators': [10, 50, 60, 200, 500, 1000],
    'learning_rate': [0.1, 0.01, 0.05],
    'eval_metric': ['rmse', 'mse', 'mape']

    grid_search = GridSearchCV(
        estimator=xg_reg,
        param_grid=parameters,
        scoring = 'neg_auc',
        n_jobs = 10,
        cv = 5,
        verbose=True
    )

    grid_search.fit(X_train, y_train)
```

Fitting 5 folds for each of 288 candidates, totalling 1440 fits

```
Out[19]: estimator=XGBRegressor(alpha=10, base_score=0.5, booster='gbtree',
                                callbacks=None, colsample_bytree=1,
                                colsample_bynode=1, colsample_bytree=0.3,
                                early_stopping_rounds=None,
                                enable_categorical=False, eval_metric=None,
                                gamma=0, gpu_id=-1, grow_policy='depthwise',
                                importance_type=None,
                                interaction_constraints='',
                                learning_rate=0.1, max_bin=256,
                                max_delta_step=0,
                                max_depth=5, max_leaves=0,
                                min_child_weight=1, missingnan,
                                monotone_constraints=(), n_estimators=10,
                                n_jobs=0, num_parallel_tree=1,
                                predictor='auto', random_state=0,
                                reg_alpha=10, ...),
        n_jobs=10,
        param_grid={'booster': ['gbtree', 'gblinear'],
                    'learning_rate': [0.1, 0.01, 0.05],
                    'max_depth': range(2, 10),
                    'n_estimators': [10, 50, 60, 200, 500, 1000]},
        verbose=True)
```

```
In [20]: grid_search.best_estimator_
grid_search.best_estimator_.fit(X_train, y_train)
y_pred_test = grid_search.best_estimator_.predict(X_test)
y_pred_train = grid_search.best_estimator_.predict(X_train)
# rmse=np.sqrt(mean_squared_error(y_test,y_pred))

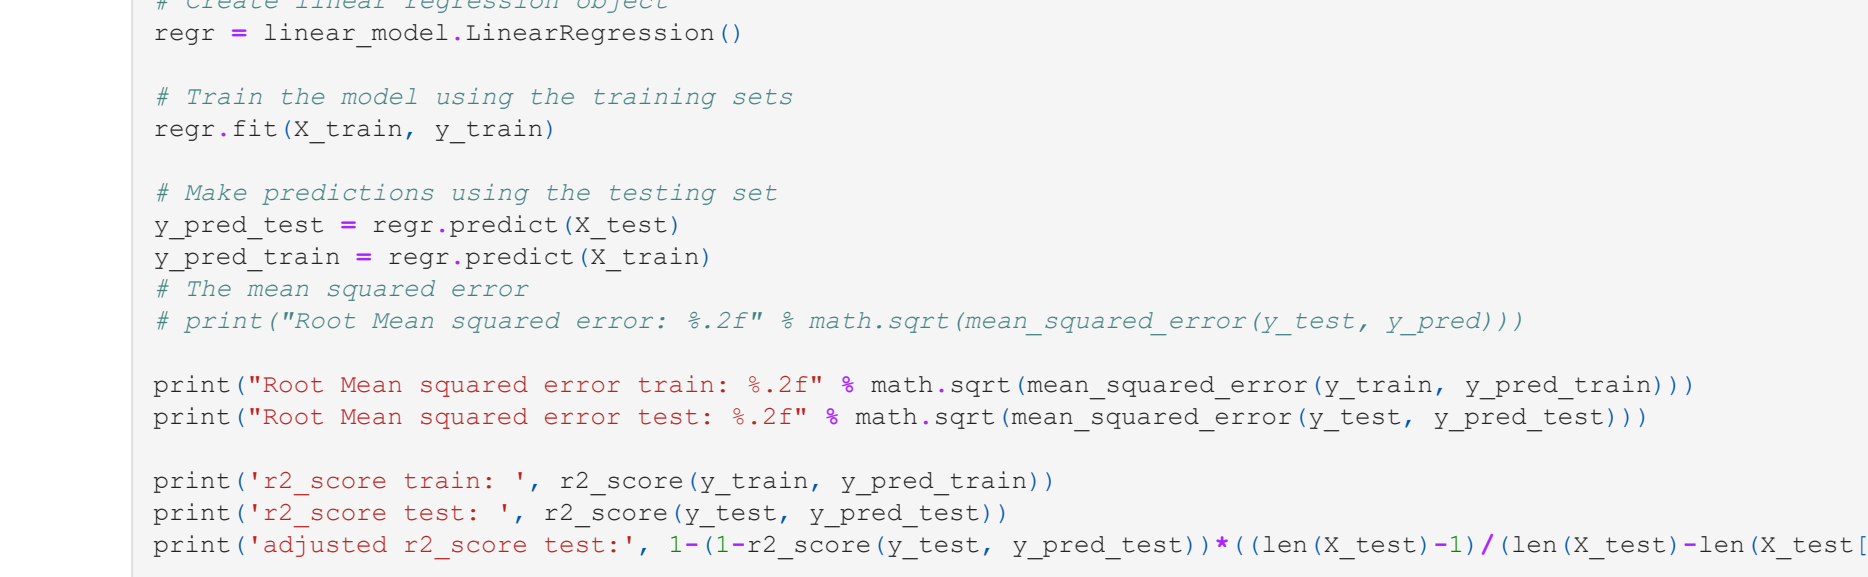

print("Root Mean squared error train: %.2f" % math.sqrt(mean_squared_error(y_train, y_pred_train)))
print("Root Mean squared error test: %.2f" % math.sqrt(mean_squared_error(y_test, y_pred_test)))

print('r2_score train: ', r2_score(y_train, y_pred_train))
print('r2_score test: ', r2_score(y_test, y_pred_test))

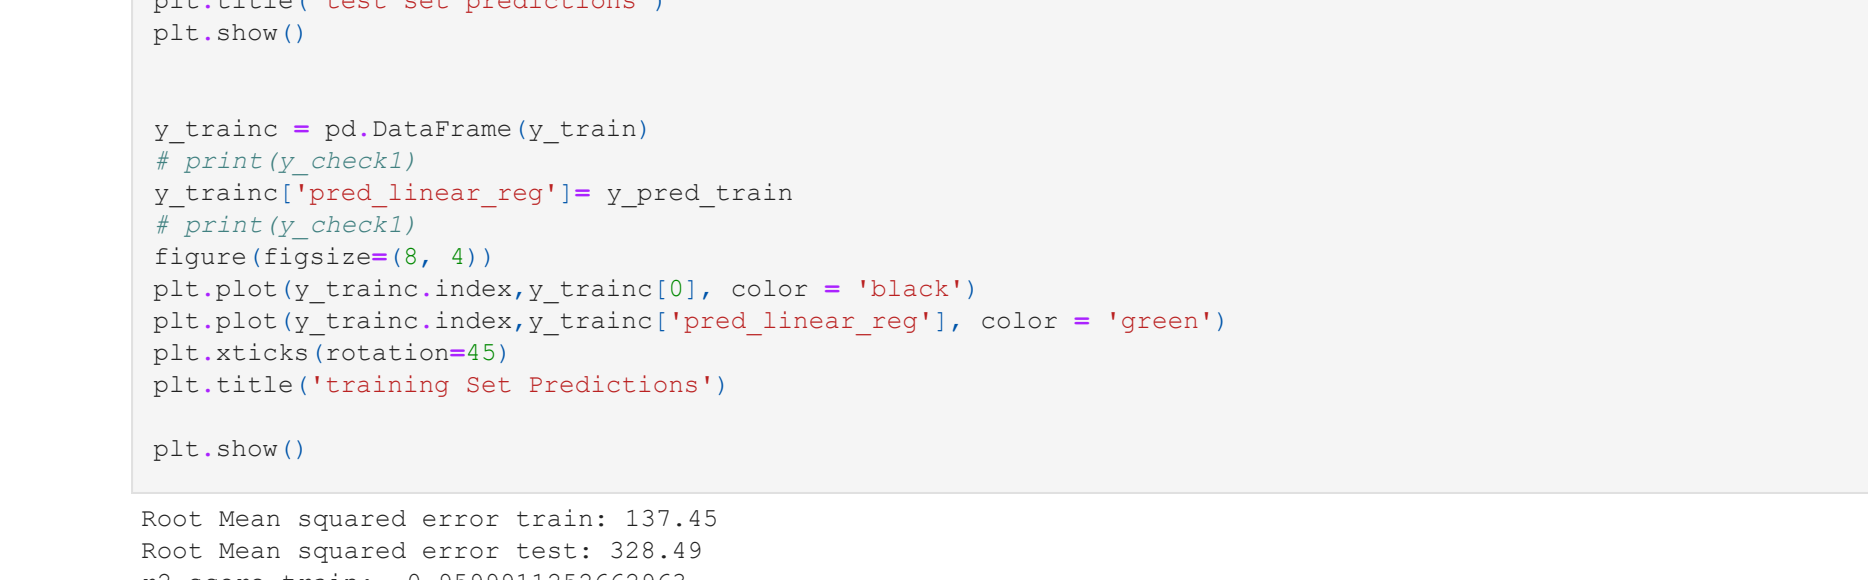
y_testc = pd.DataFrame(y_test)
# print(y_check)
y_testc['pred_linear_reg']= y_pred_test
# print(y_check)
figure(figsize=(8, 6))
plt.plot(y_testc.index,y_testc[0], color = 'black')
plt.plot(y_testc.index,y_testc['pred_linear_reg'], color = 'green')
plt.xticks(rotation=45)
plt.title('test set predictions')
plt.show()

y_trainc = pd.DataFrame(y_train)
# print(y_check)
y_trainc['pred_linear_reg']= y_pred_train
# print(y_check)
figure(figsize=(8, 6))
plt.plot(y_trainc.index,y_trainc[0], color = 'black')
plt.plot(y_trainc.index,y_trainc['pred_linear_reg'], color = 'green')
plt.xticks(rotation=45)
plt.title('training Set Predictions')
plt.show()
```

Root Mean squared error train: 117.53  
Root Mean squared error test: 695.59  
r2\_score train: 0.970746298586344  
r2\_score test: 0.43197944363017215



```
In [21]: xgb.plot_importance(grid_search.best_estimator_)
plt.rcParams['figure.figsize']= [16,10]
plt.show()
```



## linear regression on sliding window

```
In [22]: # print('linear regression with no scaling')
# Create linear regression object
reg = linear_model.LinearRegression()

# Train the model using the training sets
reg.fit(X_train, y_train)

# Make predictions using the testing set
y_pred_test = reg.predict(X_test)
y_pred_train = reg.predict(X_train)
# The Mean Squared error
# The Mean Squared error: %.2f" % math.sqrt(mean_squared_error(y_test, y_pred)))

print("Root Mean squared error train: %.2f" % math.sqrt(mean_squared_error(y_train, y_pred_train)))
print("Root Mean squared error test: %.2f" % math.sqrt(mean_squared_error(y_test, y_pred_test)))

print('r2_score train: ', r2_score(y_train, y_pred_train))
print('r2_score test: ', r2_score(y_test, y_pred_test))

print('adjusted r2_score test:', 1-(1-r2_score(y_test, y_pred_test))*((len(X_test)-1)/(len(X_test)-len(X_test)))

y_testc = pd.DataFrame(y_test)
# print(y_check)
y_testc['pred_linear_reg']= y_pred_test
# print(y_check)
figure(figsize=(8, 4))
plt.plot(y_testc.index,y_testc[0], color = 'black')
plt.plot(y_testc.index,y_testc['pred_linear_reg'], color = 'green')
plt.xticks(rotation=45)
plt.title('test set predictions')
plt.show()

y_trainc = pd.DataFrame(y_train)
# print(y_check)
y_trainc['pred_linear_reg']= y_pred_train
# print(y_check)
figure(figsize=(8, 4))
plt.plot(y_trainc.index,y_trainc[0], color = 'black')
plt.plot(y_trainc.index,y_trainc['pred_linear_reg'], color = 'green')
plt.xticks(rotation=45)
plt.title('training Set Predictions')
plt.show()
```

Root Mean squared error train: 137.45  
Root Mean squared error test: 328.49  
r2\_score train: 0.959891252662983  
r2\_score test: 0.8733295124304821  
adjusted r2\_score test: 0.8700634501786226

